



Two builders who kept seeing the same gap, and turned co-founders.



Topaz

Product and go-to-market. Just finished the MIT Sloan MBA.



Technion

MIT

intel

IBM

aws



Omer

Security and architecture. From the Intelligence Corps into cybersecurity.



Intelligence Corps



Data Architect



Cybersecurity

We worked together for years in Tel Aviv, then reconnected as **co-founders** to build a secure **AI agents platform for enterprises**. We came to this hackathon as **mentors**. On day two we decided to build for the developer community too. **48 hours later**, this is IronClaw.

• SECURE · SELF-HOSTED · OPEN SOURCE



IRONCLAW

AI agents you can **actually trust**.

A secure developer control plane for OpenClaw-style agents.
Sandboxed, governed, and auditable from the first line.

Written in **Go**

gVisor isolation

Encrypted queues

MIT · open source

• **BUILT IN 48 HOURS**



The more useful an agent gets, the more dangerous it gets.

An agent only earns its keep if it can act on its own. But every action it can take is also a way to do harm, and today those actions land in the most sensitive parts of your stack:

+ your code repos

+ the terminal

+ secrets and tokens

+ Slack · Discord · Telegram

+ browsers and APIs

+ long-running memory

The usual reassurance is "we keep audit logs." But an agent that can rewrite its own code can rewrite how it logs. **Auditing you cannot trust is not auditing.**



Adoption is here. Trust is the bottleneck.

Developers are not waiting to be convinced that agents are useful. They already use them. What they do not have yet is a reason to trust them with real access.

\$307B → \$632B

Projected AI spend, 2025 to 2028. Budgets are big enough to fund a new infrastructure layer.

SOURCE **IDC**

84%

of developers use or plan to use AI tools. The adoption funnel already exists.

SOURCE  **Stack Overflow**

1.1M+

public repos already wiring an LLM SDK into real code, up sharply year over year.

SOURCE  **GitHub**

Existing tools prove the demand. **The security layer is the piece nobody owns yet.**



A secure developer **control plane** for agents.

We give developers OpenClaw-style power, but we treat the agent as untrusted from the first line. The boundary is something you can verify, not something you take on faith.



Sealed runtime

Compiled Go. The agent **cannot rewrite itself**, so it cannot forge its own audit trail.



Sealed sandbox

gVisor, **network=none**. The model is reachable only through a host proxy.



Encrypted queues

Per-session keys. Inbound is **read-only by design**. No cross-session reads.



Approval gateway

Only the agent's **own capability changes** stop for a human. It works freely otherwise.

The agent does its job on its own, but it can **never grant itself new powers**, and it cannot tamper with the record. The trust sits outside the agent, where it cannot reach.



The market already proves the demand.

PROJECT	TRACTION SIGNAL	WHAT IT PROVES	THE GAP WE FILL
OpenClaw	~379K ★ · 104K Reddit · 448K video views	Breadth and mindshare	Security is not the product story
Hermes	~196K ★ · 124K Reddit	Rich capability, learning loop	Broad surface means more risk
NanoClaw	~30K ★ · per-agent containers	Isolation matters	Still early and narrow
LangGraph	~35K ★ · orchestration standard	Agent orchestration	Not a secure execution product
OpenHands	~77K ★ · coding-agent workflows	Dev workflows are real	Security depends on how you deploy
IronClaw	Secure by default · open source · Go + gVisor	Trust is the product	The one a security team can approve

OpenClaw proved demand. NanoClaw proved isolation matters. IronClaw makes security the default.



Start with developers. **Expand to teams.**

\$52_B

TAM · the agentic-AI market by 2030, growing **46% a year**. Every autonomous agent needs a secure runtime to act in the real world. That is the layer we own.

\$9.5_B

SAM · the secure-runtime slice: 47M developers × 84% adopting agents × \$20/mo, the orgs that will not run agents without isolation and audit.

\$38_M

SOM · 3-year ARR target. Open-source wedge → enterprise governance, seeded by 20 design partners, about 0.4% of SAM.

Developer wedge **PLG**

Open source, one-command setup, CLI and chat first, secure defaults out of the box. Useful in five minutes.

expands into ↓

Enterprise **GOVERNANCE**

SSO, RBAC, policy-as-code, audit export, central secret management, deployment controls.

One funnel, one unit: a **\$52B** agent market, the **\$9.5B** slice that demands isolation and audit, and the **\$38M** we can realistically win in three years. We do not compete with Cursor or Copilot. We are the secured layer their agents run inside. Those seat prices already prove the market pays for agent tooling; we price the control plane beneath it.

Win the developer in **five minutes**. Win the company by making **security review easy**.



It is **real, runnable**, and on GitHub.

ironctl · live on a sealed, sandboxed runtime

```
# one command installs the two host binaries
$ curl -fsSL ../install.sh | sh    ✓ checksum verified
$ ironclaw-controlplane --dev    ✓ api on 127.0.0.1:8787

# the agent asks to change its own powers
$ ironctl change submit --kind persona --group default
  ↳ HELD at gateway, waiting for a human
$ ironctl change pending          1 change waiting
$ ironctl change approve <id>    ✓ applied · written to audit log
```

ENFORCED BY DESIGN, NOT BY GOOD BEHAVIOR

- ✓ **Sealed Go binary**: no source in the box to rewrite, so it cannot fake its own audit trail.
- ✓ **gVisor, network=none**: even a hostile agent has nowhere to send your data.
- ✓ **Per-session encrypted SQLite**: keys are host-held; no session can read another.
- ✓ **Capability changes can't self-apply**: the agent can ask for power, never take it.

Go + CGO

gVisor network=none

Encrypted SQLite queues

~200 tests green

CI passing

6 channel adapters

MIT licensed

Channels shipped: webhook, Telegram, Slack, Discord, Matrix, Google Chat. Two compiled binaries that never share memory and talk only through encrypted per-session files.



Now help it **get everywhere.**

IronClaw is open source, it runs today, and it is the secure foundation the agent wave is missing. We are not here to raise money, we are here to get it into the right hands. Here is how each of you can help:



ClawComp

Showcase it. Feature IronClaw as something real that was built here, and put it in front of the next competitions.



Builder & AI-first teams

Adopt it. Pull the IronClaw suite into your products, so your agents get real access and ship secure from day one, not patched later.



Investors

Recommend it. If you want your portfolio to go AI-first without going exposed, send them the project.



Developers

Contribute to it. It is open source. Send a PR, harden a wall, and make IronClaw's security stronger for everyone who builds on it.

Open source. Running today.

IronClaw is the next big thing in secure agents, so help us put it everywhere.